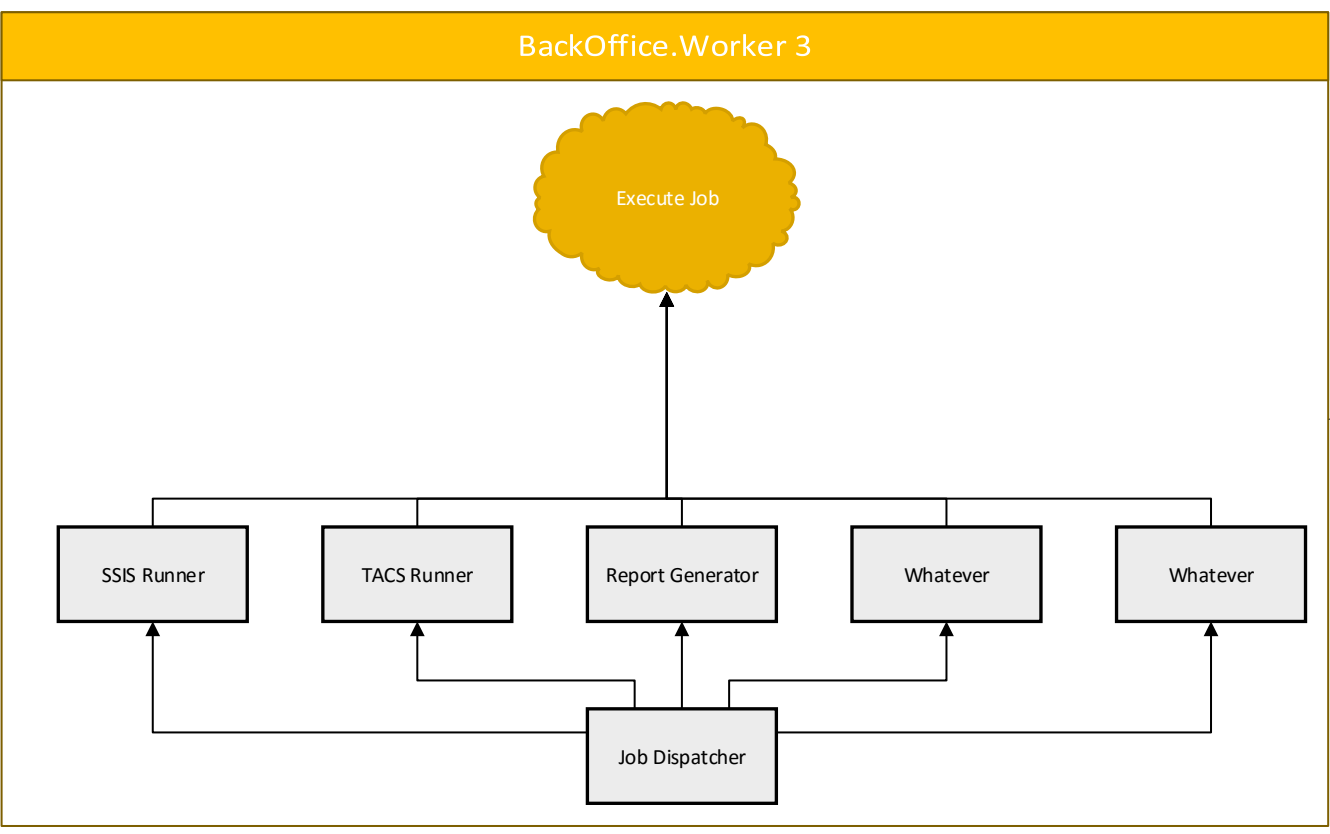
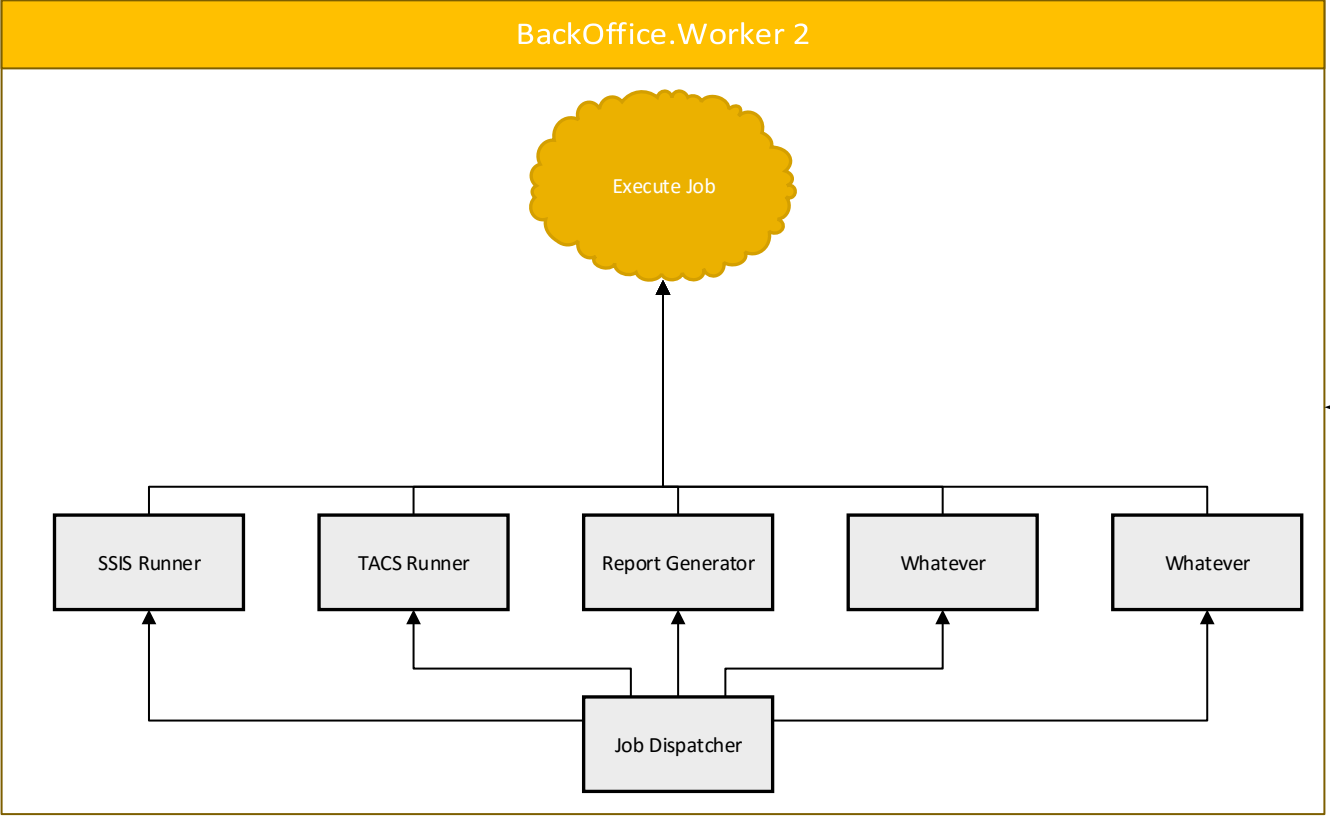
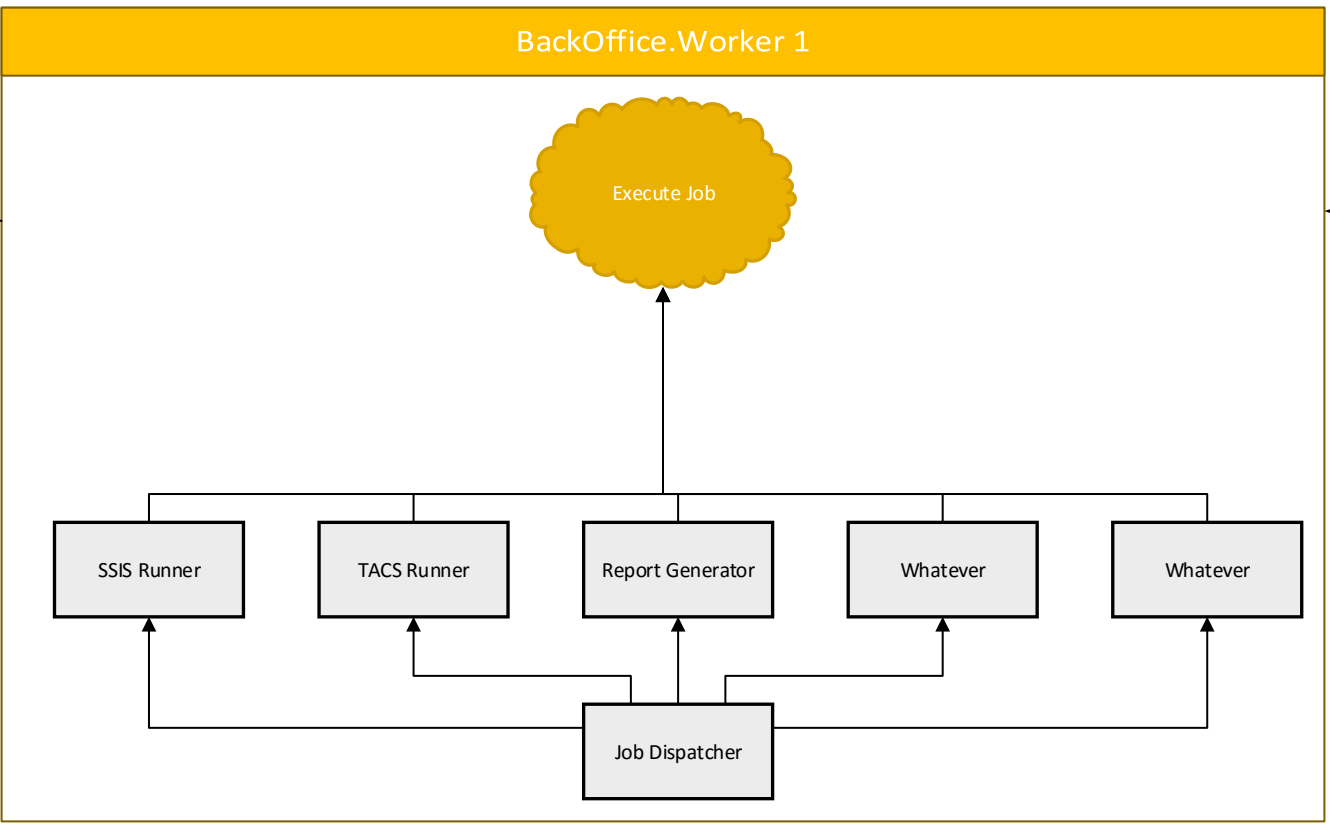


- Scalable
- Can execute any job
- Runners as dependencies to reduce complexity
- The only component that needs to be frequently released
- Reports job progress
- In case of job failure, do not removes job from the queue or puts it into special queue



We can use any other queue but MongoDB seems to be the best option because:

- it is free
- easily scalable
- supports atomic read/update messages without pulling them from queue (RabbitMQ does not supports this)
- perfect for json data

